

PRIYA SHARMA

Email: priya.sharma@mleng.dev | Phone: +91 9988776655
LinkedIn: linkedin.com/in/priyasharma-ml | GitHub: github.com/priyasharma
Location: Bangalore, India

PROFESSIONAL SUMMARY

Machine Learning Engineer with 3+ years of experience specializing in scalable ML systems, MLOps, deep learning architectures, and production model serving on Kubernetes and AWS.

EDUCATION

B.Tech in Computer Science & Engineering (2019 - 2023)
VIT Vellore | CGPA: 8.9 / 10.0

TECHNICAL SKILLS

- Languages: Python, C++, SQL, Bash
- Frameworks & DL: PyTorch, TensorFlow, Scikit-learn, ONNX, TensorRT
- MLOps & Infrastructure: Docker, Kubernetes, CI/CD, MLflow, Kubeflow, AWS (EC2, S3, SageMaker, Lambda), Linux
- Backend & Serving: FastAPI, gRPC, Triton Inference Server, Redis, PostgreSQL
- Engineering Practices: System Design, Git, Agile, Unit Testing, Distributed Training

EXPERIENCE

Machine Learning Engineer | Apex AI Systems (June 2023 - Present)

- Designed and maintained automated CI/CD MLOps pipelines using Docker, GitHub Actions, and Kubernetes, deploying 12+ computer vision and NLP models.
- Optimized PyTorch model inference using ONNX Runtime and TensorRT, achieving 3.8x throughput speedup and cutting cloud GPU compute costs by 45%.
- Implemented real-time data drift and concept drift monitoring with Prometheus and MLflow.

PROJECTS

1. Real-time Distributed Model Inference Engine

- Architected a high-concurrency model serving platform using FastAPI, Triton, and Redis caching handling 1,500 req/sec at <30ms latency.
- Tech Stack: Python, PyTorch, Docker, Kubernetes, Triton, Redis, AWS.

2. MLOps Automated Training & Deployment Pipeline

- Built end-to-end reproducible pipeline from S3 data ingestion to automated model evaluation and containerized deployment.
- Tech Stack: MLflow, Docker, AWS SageMaker, GitHub Actions, PyTorch.